

Computer vision: project

Thomas Aelbrecht, Andreas De Witte, Jochen Laroy, Pieter-Jan Philips, and
Gillis Werrebrouck

Ghent University, Valentin Vaerwyckweg 1, 9000 Ghent, Belgium

Abstract. The abstract should briefly summarize the contents of the paper in 150–250 words.

Keywords: Painting detection · Painting matching · Hidden Markove model.

Introduction

This paper focusses on describing the different steps to reconstruct the path of an excursion through a museum. The first section describes how the paintings are cut out of high quality pictures. The following section explains an algorithm that can independently find paintings in a picture. This section also describes how accuracy metrics can be collected, which is useful to search for possible improvements. The third section is about how paintings are being matched. This section describes performant solutions to figure out which paintings are visible in a picture. The fourth section is about how to localize where a person is. This is done by using a custom implementation of the Hidden Markov model to decide which location is a logical result. The fifth and final section is about the visualization of the result. This section describes a modern way of visualizing the visited rooms in the museum.

1 Semi-supervised painting detection

1.1 The naive painting detection algorithm

The contour detection is the core part of the semi-supervised painting detection. It's the part that decides what contours are in the image and it works as described below.

Before any contour detection can be done, as many unnecessary details as possible have to be removed. OpenCV has several options to solve this problem, for example eroding, dilating, blurring (median, Gaussian...) or downscaling. The image needs to get rid of as many details as possible to prevent detection of contours inside the paintings or on the wall. Although this won't be perfect, the removal of details will decrease the number of incorrectly detected contours.

So the first step of the algorithm is to remove details by resizing the image. The current implementation scales the image down and back up with a factor

5. To scale the image back up, pixel values are calculated using a pixel's area. The scale up was mostly done to be able to show the original image, but it's also a possibility to use the algorithm with the downscaled image. This will result in a slight increase in performance because the image it's working on would be smaller.

The next step is to convert the image to grayscale. This is done to make it easier to differentiate certain parts. This grayscaled image is then dilated and eroded to remove even more noise at the borders of the painting or on the walls. The last step of removing noise is applying a median blur. The biggest advantage of using a median blur is that it will preserve edges while removing noise. The idea behind these steps is to smeaere as many details as possible, in other words to make bigger blots with the same color. This makes it easier to detect the borders and remove noise in the background.

The next step is to detect edges with Canny. The result of the Canny function will then be dilated once again to make the found edges stronger. Thereafter the contours can be detected using OpenCV's contour detection algorithm. This algorithm will make sure only the most outer contours are returned when a hierarchical structure of contours is found. A small final detail to prevent unlogical solutions is the following: any contour with a ratio smaller than 1:10 will be removed. This prevents very small contours to appear around noisy parts of the image.

1.2 Strengths and weaknesses

This biggest strength of this algorithm is that it will give an output in almost every image. An example of an image where no painting will be detected is an image where the framework of the painting is invisible, meaning the image only contains the painting itself, no wall and no framework. With this kind of paintings, the desired solution is a contour containing the entire painting, while this is impossible because the painting has no border at all. Another advantage is that the algorithm is very fast in detecting contours due to the fact that all parts of the algorithm are standard functions that have a good performance.

However, the biggest weakness of this algorithm is that almost every found contour is not precise enough, so almost every solution needs a slight adjustment. For some solutions, the algorithm detects contours that are slightly larger than the actual painting, while for other paintings, it doesn't even include the border of the painting.

Another weakness of the algorithm is that paintings sometimes have overexposure or shadows, which makes it harder to detect it's contours. Some paintings even have a small tag with a description next to the painting, sometimes this small tag is detected as being a painting. This is a logical decision, because there's a big contrast in colors and the tag has a clear contour, but this is not a desired effect.

This algorithm is a first version to quickly be able to fill the database with the groundtruth. The actual algorithm in its final version is completely different

and doesn't have these issues anymore. This algorithm will be discussed in the next section.

2 Unsupervised painting detection

2.1 The detection algorithm

The second detection algorithm drastically differs from the first attempt (see Section 1.1), it was designed from scratch with the weaknesses of the previous algorithm in mind.

The first step in the algorithm is a mean shift segmentation (see Figure 1)). This is a method to remove noise by taking the mean of the pixels within a certain range. A big advantage is that it partially removes color gradients and fine-grain textures which would cause issues in the further steps of the algorithm.



Fig. 1. Mean shift segmentation.

A common technique in object detection is by creating a mask of the object. The problem with this is that there's no certainty about the size and ratio of the painting as well as where the painting is located in the image and how many there are in the image. This can be solved by thinking the other way around. The one certainty that is consistent throughout all images is that all paintings hang on a wall. A mask for the wall will be made instead of making a mask for the paintings. The mask of the painting(s) can then be obtained by simply inverting this mask.

Another issue is that this mask can't be statically programmed in code for each image since the algorithm needs to be usable on any image or video frame. This can be solved by using a technique called flooding. Flooding will create a mask from a starting position in the image and will fill all neighbouring pixels

if they have a color close to the color of the starting position. This is done recursively until no pixels are within the color range of the starting position anymore.

add reference to links underneath in comment




Fig. 2. Wall mask.



Fig. 3. Paintings mask.

The next problem is that there needs to be a good starting position to have a good and correct mask for the wall. To find this, the unsupervised algorithm will iterate over all pixels of the image with a certain step size and perform the flooding with that pixel as starting position. Once a mask that has the same height and width as the image is found, then this mask will be used in the next steps of the algorithm. If no such mask is found after iterating over all pixels (with a certain step size), then the mask with the biggest size is used. This is because sometimes parts of the floor or other elements in the image will obstruct the floodfill algorithm to find a mask that has the same size as the image. See Figure 2) for an example of such mask.

Once the mask of the wall has been obtained, it is inverted to have the mask of the paintings (see Figure 3). The mask is then eroded to remove small imperfections and a median blur is used to smooth the edges in the mask.

The next step is to use Canny, the difference with the naive approach is that the two threshold values are determined by using the Otsu algorithm.

Next, a morphological transformations is performed on the Canny edges (see Figure 4). This will close edges that are close to each other as this will improve the detection of closed contours. This morphological transformation is in essence a dilation followed by an erosion.

The next steps are exactly the same as the naive painting detection algorithm as described before. The contours are detected in the Canny edges and only fully enclosed polygons with 4 sides are returned as quadrilaterals.




Fig. 4. Paintings edges.

2.2 Strengths and weaknesses

The most important improvement of the detection algorithm is that it isn't based on smearing out colors and finding contours in the color smudges.

The key to success with this algorithm is the use of automatic mask creation by using floodfill as main technique. Most of the paintings are detected in most of the images. The detected paintings also have a better fit compared to the naive algorithm.

Most paintings in the dataset are correctly detected as well those in the test set. Only on some occasions there is a miss detection. One flaw that has been found is that some doorways are detected as paintings. This is logical because of how the mask creation works with the floodfilling. In the case of such a detection flaw, the doorways obstructs the floodfilling algorithm to correctly create a mask of the non-painting area. However, this is something that doesn't happen too often.

A minor problem with this algorithm, which is also the case in the naive algorithm, is that it sometimes detects dark shadows as part of the painting. Although, this is less of an issue in this algorithm than it was in the naive algorithm. It doesn't affect the accuracy too much because if there is a shadow, it only appears on one side of the painting and it doesn't reach far.

A drawback of this algorithm is that it is less performant than the naive algorithm. This is mainly because of the mask creation, specifically the floodfilling. To solve this problem the critical and less performant parts of the algorithm were transformed into Cython. This resolved the whole performance issue and increased the processing performance drastically.

add reference about cython to links underneath in comment

2.3 Quantitative comparison

In order to make a quantitative comparison, two things are needed. First of all, the quadrilaterals have to be found autonomously (as described in Section 1.1). The second thing needed is the groundtruth of the dataset which has been generated by using the solution created with the first algorithm (see Section 1).


To measure the accuracy of the painting detection algorithm, three things are required; the amount of false negatives (= paintings that aren't found at all), the amount of false positives (= detected paintings that aren't paintings) and the bounding box accuracy (= average intersection divided by union).

With the creation of the solution for this problem, a new problem occurred: how to find the intersection of these two shapes? The solution to this problem is made by using "Shapely". To do so, the quadrilaterals have to be transformed into a polygon. Once this is done, "Shapely" can compute the intersection. It has been tested whether this gives correct intersections if two polygons are not intersecting, or when they are sharing only a line. Also some more general intersections have been tested. Once the intersection is made, it's easy to find the area of the polygon, using "Shapely" once again.


The accuracy results for this detection algorithm were fairly good. An example of the detection can be found in Figure 5. The detected contours are shown in red and the groundtruth is shown in blue as determined with the use of the naive algorithm. The dataset consists of 553 images which all together contain 836 paintings. The detection algorithm had 86 false negatives (paintings) and 32 false positives (paintings). The average bounding box accuracy is 81.30%. These results are displayed on the graphs in Figure 6.




Fig. 5. Paintings detection.



(a) Average room detection accuracy



(b) Detection accuracy rangesc for the dataset



(c) False negatives and false positives per room

Fig. 6. Performance measurements

2.4 Qualitative evaluation

For the qualitative evaluation, the unsupervised painting detection algorithm is used to check how accurate the paintings are found on a more difficult set. Hereby the test set and the video files are used. The painting detection is working well, most paintings are being found. However, some difficulties occur. The algorithm sometimes defines a window as being a painting and it almost always defines a doorway as painting.

Another problem is images shot at a sharp angle. For example when it seems like the borders of to adjacent paintings are overlapping. These images are sometimes not recognized. Images that are too blurred are sometimes not found either.

3 Matching

3.1 Features

Before diving deeper into the matching algorithm, it's important to explain the features that used to do the matching. The algorithm only depends on two types histograms of the detected paintings. The first type is a histogram of the full painting, the second type is a collection of histograms gathered from different blocks of the painting. The block sized used in this algorithm divides the painting in 4 rows by 4 columns, independent of the painting's size. Both features are saved to the database for each of the labeled paintings and are both used in the matching algorithm explained in Section 3.2.

3.2 The matching algorithm

The first step in matching a given painting with the entire dataset of paintings is a light intensity equalization. This way different light intensities have no influence on the histograms gathered from the detected painting. Equalization is a technique that is only used on the video frames because of the rather bad frame quality, the dataset itself and the previously performed benchmarks don't need this step.

add reference about equalization to links underneath in comment

The second step consists of fetching the histograms as described in Section 3.1. Thereafter these histograms are compared against all histograms in the dataset. Per known painting the distance between the histograms is calculated using a correlation. This way a chance between 0 and 1 is obtained from this calculation. The histogram of the whole painting gives one chance, the block histogram gives a total of 16 chances which are combined by taking the average. These two chances are combined using Formula 1 which calculates a weighted average. In this formula $P(X = P_i)$ stands for the chance that the detected painting X is painting P_i from the dataset, B stands for the average of the block histogram distances and F stands for the distance of the normal histograms.

The block histogram gets a much higher weight because the predictions of these histograms are much more reliable than these of the normal histograms.

$$P(X = P_i) = \frac{(8 * B) + (1 * F)}{9} \quad (1)$$

The matching algorithm gives per detected painting a list of possible rooms with the according chances, even if the room is not possible. How impossible rooms are treated, will be explained in Section 4. It's also possible to ignore chances that are below a given threshold.

3.3 Matching results

The matching algorithm has been tested in three ways. First, the dataset was matched to itself without using the detection algorithm but by using the corners from the semi-supervised algorithm. The result for this test is that 835 of 836 paintings were matched correctly. This means that only one painting could not be uniquely matched, this could be because some paintings appear in multiple images. The average matching probability is 92.24%.

The second test was on the dataset but with the detection algorithm. The result for this is that 702 of the 762 detected paintings were matched correctly.

The third and last test was on the test dataset also with the detection algorithm. The images have been divided in folders with the correct room as name to be able to count the amount of correct room matches. The result for this test is that 158 of the 338 detect paintings had a correct room match. This is lower than the dataset because these images are meant to be more difficult to match because of the angle, rotation, light intensity, etc. These results of prediction on the dataset are displayed on the graphs in Figure 7.




Fig. 7. Correct and incorrect painting predictions per room

4 Localization

4.1 Hidden Markov model

In order to establish an intelligent localization of the user by eliminating teleportations, a model inspired by the Hidden Markov model [1], was designed. The reason why no pure Hidden Markov can be used is because the input and output of the model is identical. Another reason is that there's no way to determine the initial room of a given video without having a number of observations, which are rooms in this case. Therefor an alternative Hidden Markov model was designed.

add reference to link underneath

4.2 Alternative Hidden Markov model

The alternative model basicly keeps a history of n observations and predicts the most common observation as the current room. Initially the model doesn't know the current room and therefor it doesn't emit any prediction until enough observations are obtained.

The model takes a list of probabilities per painting as an input which could possibly contain different probabilities for the same room. However the model needs one probability per room. So Formula 2 [2] is used to combine all probabilities for a given room into one probability. In this formula, $P(X = R)$ stands for the chance that the user is in room R , $P_{test_i}(X = R)$ stands for the chance given by $test_i$ that the user is in room R and $P_{test_i}(X \neq R)$ is therefor the chance that the user is not in room R .

$$P(X = R) = \frac{\prod_i P_{test_i}(X = R)}{\prod_i P_{test_i}(X = R) + \prod_i P_{test_i}(X \neq R)} \quad (2)$$

Although this formula can combine probabilities, it doesn't take the current room into account. Therefor a weight is added to each of the chances. This weight is equal to 1 when the room is accessible from the current room and can be choosen for rooms that aren't accessible, the default is 0.5. This results in Formula 3.

$$P(X = R) = \frac{\prod_i w * P_{test_i}(X = R)}{\prod_i w * P_{test_i}(X = R) + \prod_i w * P_{test_i}(X \neq R)} \quad (3)$$

By calculating the result of this formula per room, a list of probabilities per room is obtained. This list is then used to find the room with the highest probability which will be the observation for this input.

If the current observation is accessible from the current room, it is added to the history of observations. If this is not the case, the current room is added the history. Now, the most common room in the history is taken as the prediction for this input.

5 Visualization

In order to create a visualization of the algorithms described in the previous sections, a 3D model of the floor plan of the museum was created. First, this floor plan is converted to an SVG image so it can easily be altered in Python. Working this way makes it possible to dynamically create an HTML page showing the SVG floor plan and the currently analyzed frame along with the detected paintings. Because some operations need a lot of processing power, two separate processes are spawned. One process handles the painting detection and room prediction, the other one handles the visualization of the HTML page in a webview. One of the major points of attention during the creation of the visualization is performance, hence the HTML file is never written to disk in order to bypass the IO operations. Therefore interprocess communication is used to send the HTML page to the second process. When this process receives an HTML page, it can be visualized in the webview. An example of this webview is shown in Figure 5. Besides showing the floor plan and the image, other information is shown as well. This includes information about the predicted room, which is also visualized on the floor plan, a probability for this prediction and the currently processed video's name.




Fig. 8. Modern visualization of the floor plan

References

1. Eddy, S.R.: Hidden markov models. *Current opinion in structural biology* **6**(3), 361–365 (1996)
2. Genest, C., Zidek, J.V., et al.: Combining probability distributions: A critique and an annotated bibliography. *Statistical Science* **1**(1), 114–135 (1986)